

ouees-202006 topic 10:

Network fault-tolerance

Kenji Rikitake

10-JUN-2020

School of Engineering Science, Osaka University

On the internet

@jj1bdx

Copyright ©2018-2020 Kenji Rikitake.

This work is licensed under a [Creative Commons Attribution 4.0 International License](#).

CAUTION

Osaka University School of Engineering Science prohibits copying/redistribution of the lecture series video/audio files used in this lecture series.

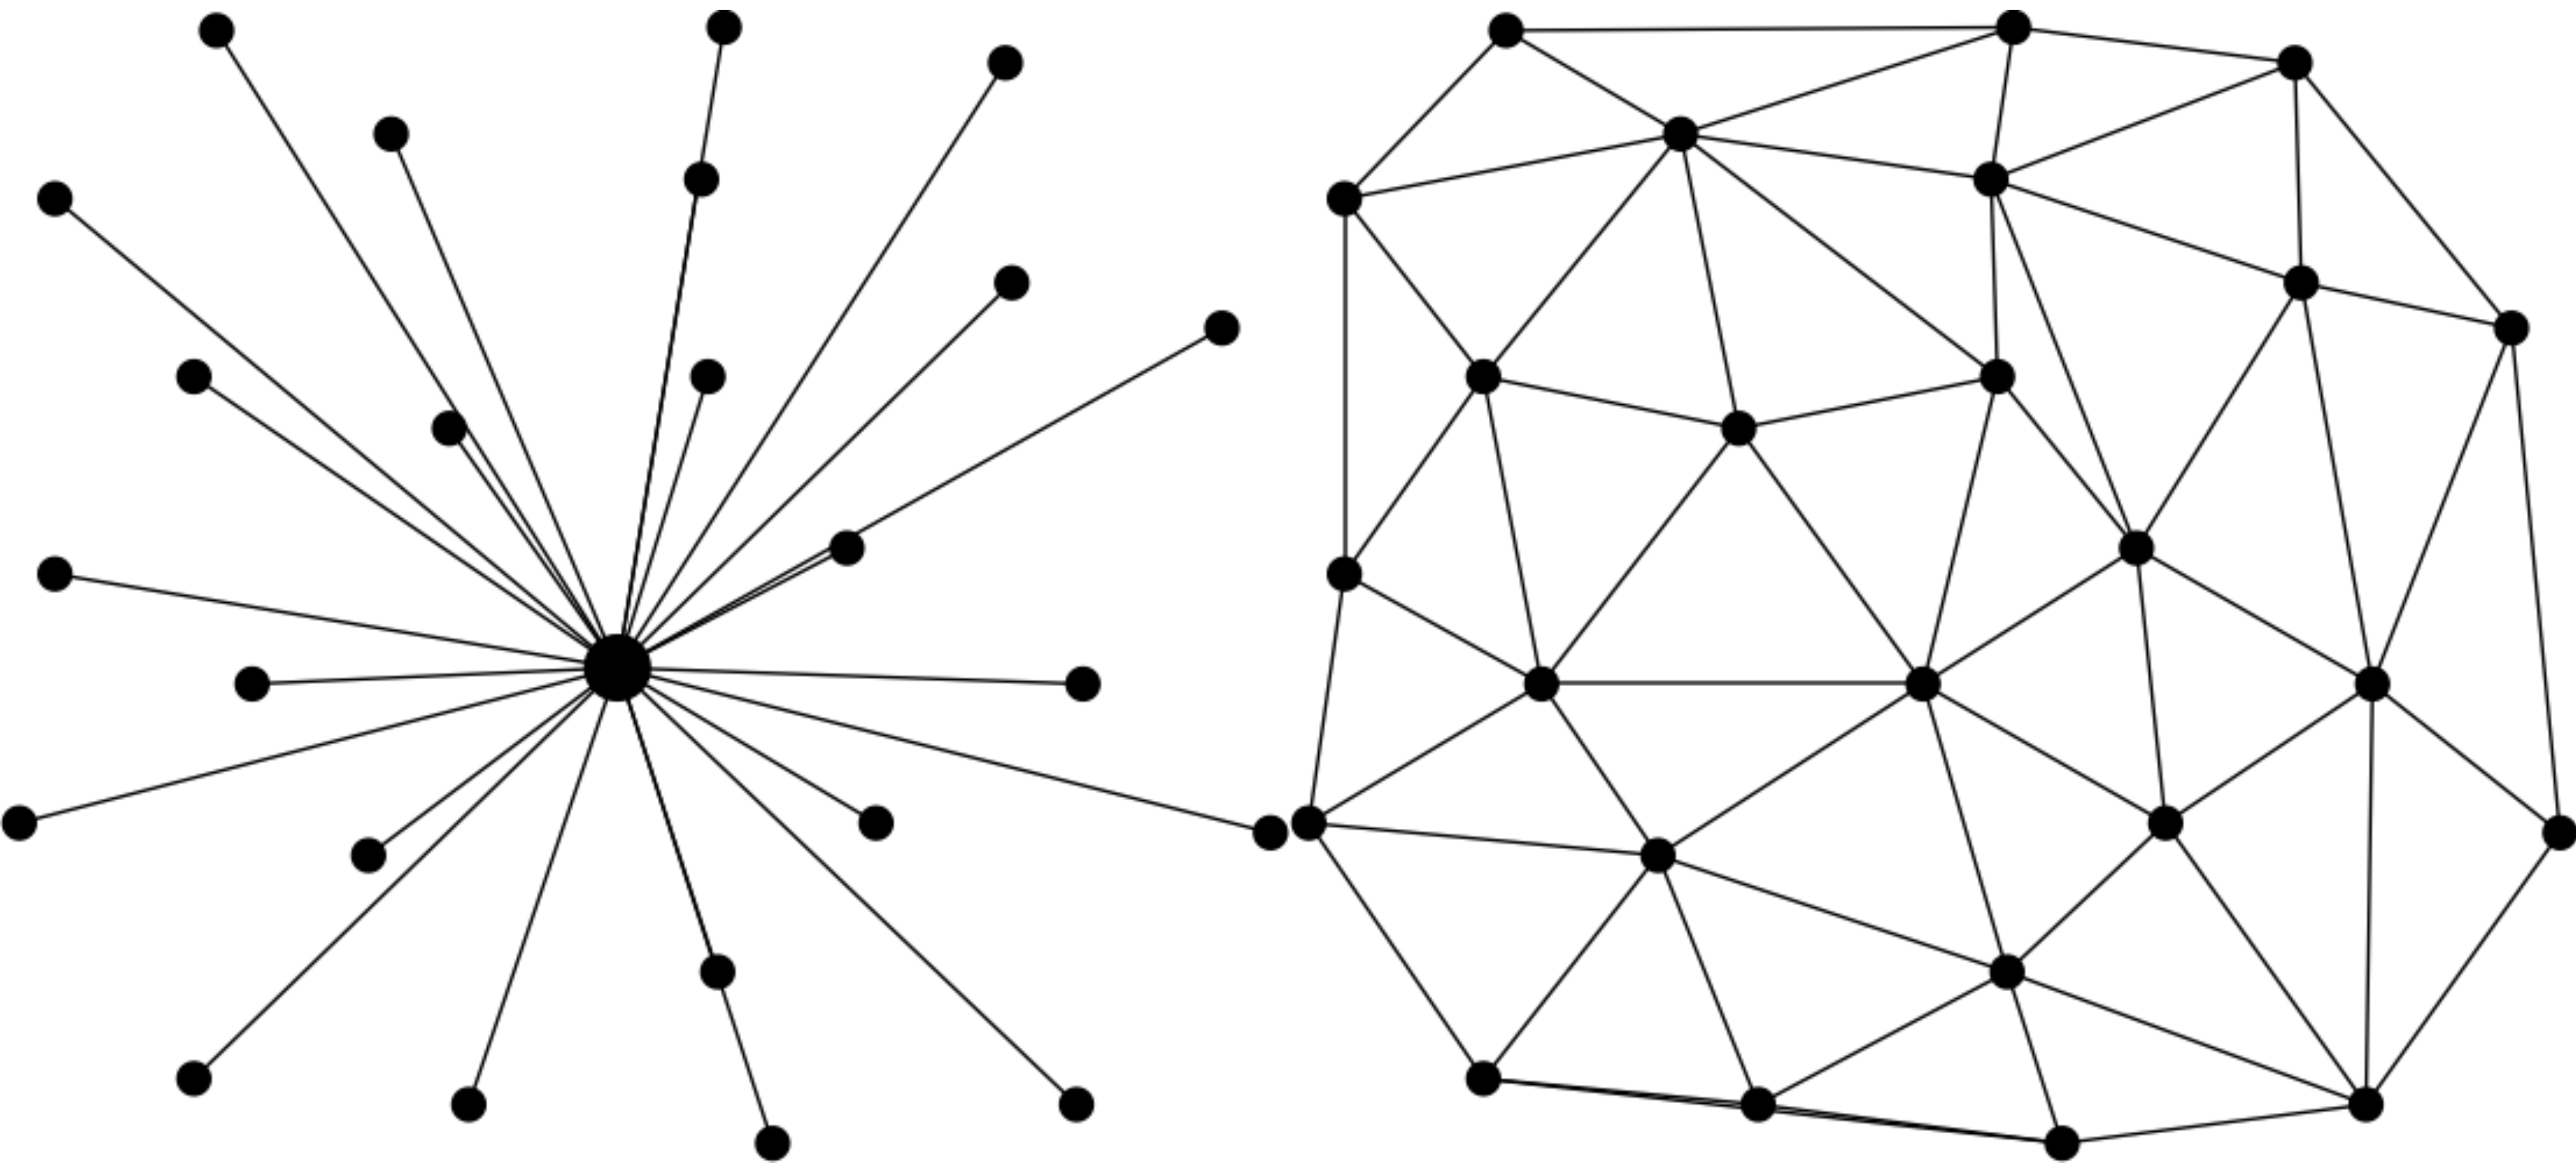
大阪大学基礎工学部からの要請により、本講義で使用するビデオ/音声ファイルの複製や再配布は禁止されています。

Lecture notes and reporting

- <https://github.com/jj1bdx/oueees-202006-public/>
- Check out the README.md file and the issues!
- Keyword at the end of the talk
- URL for submitting the report at the end of the talk

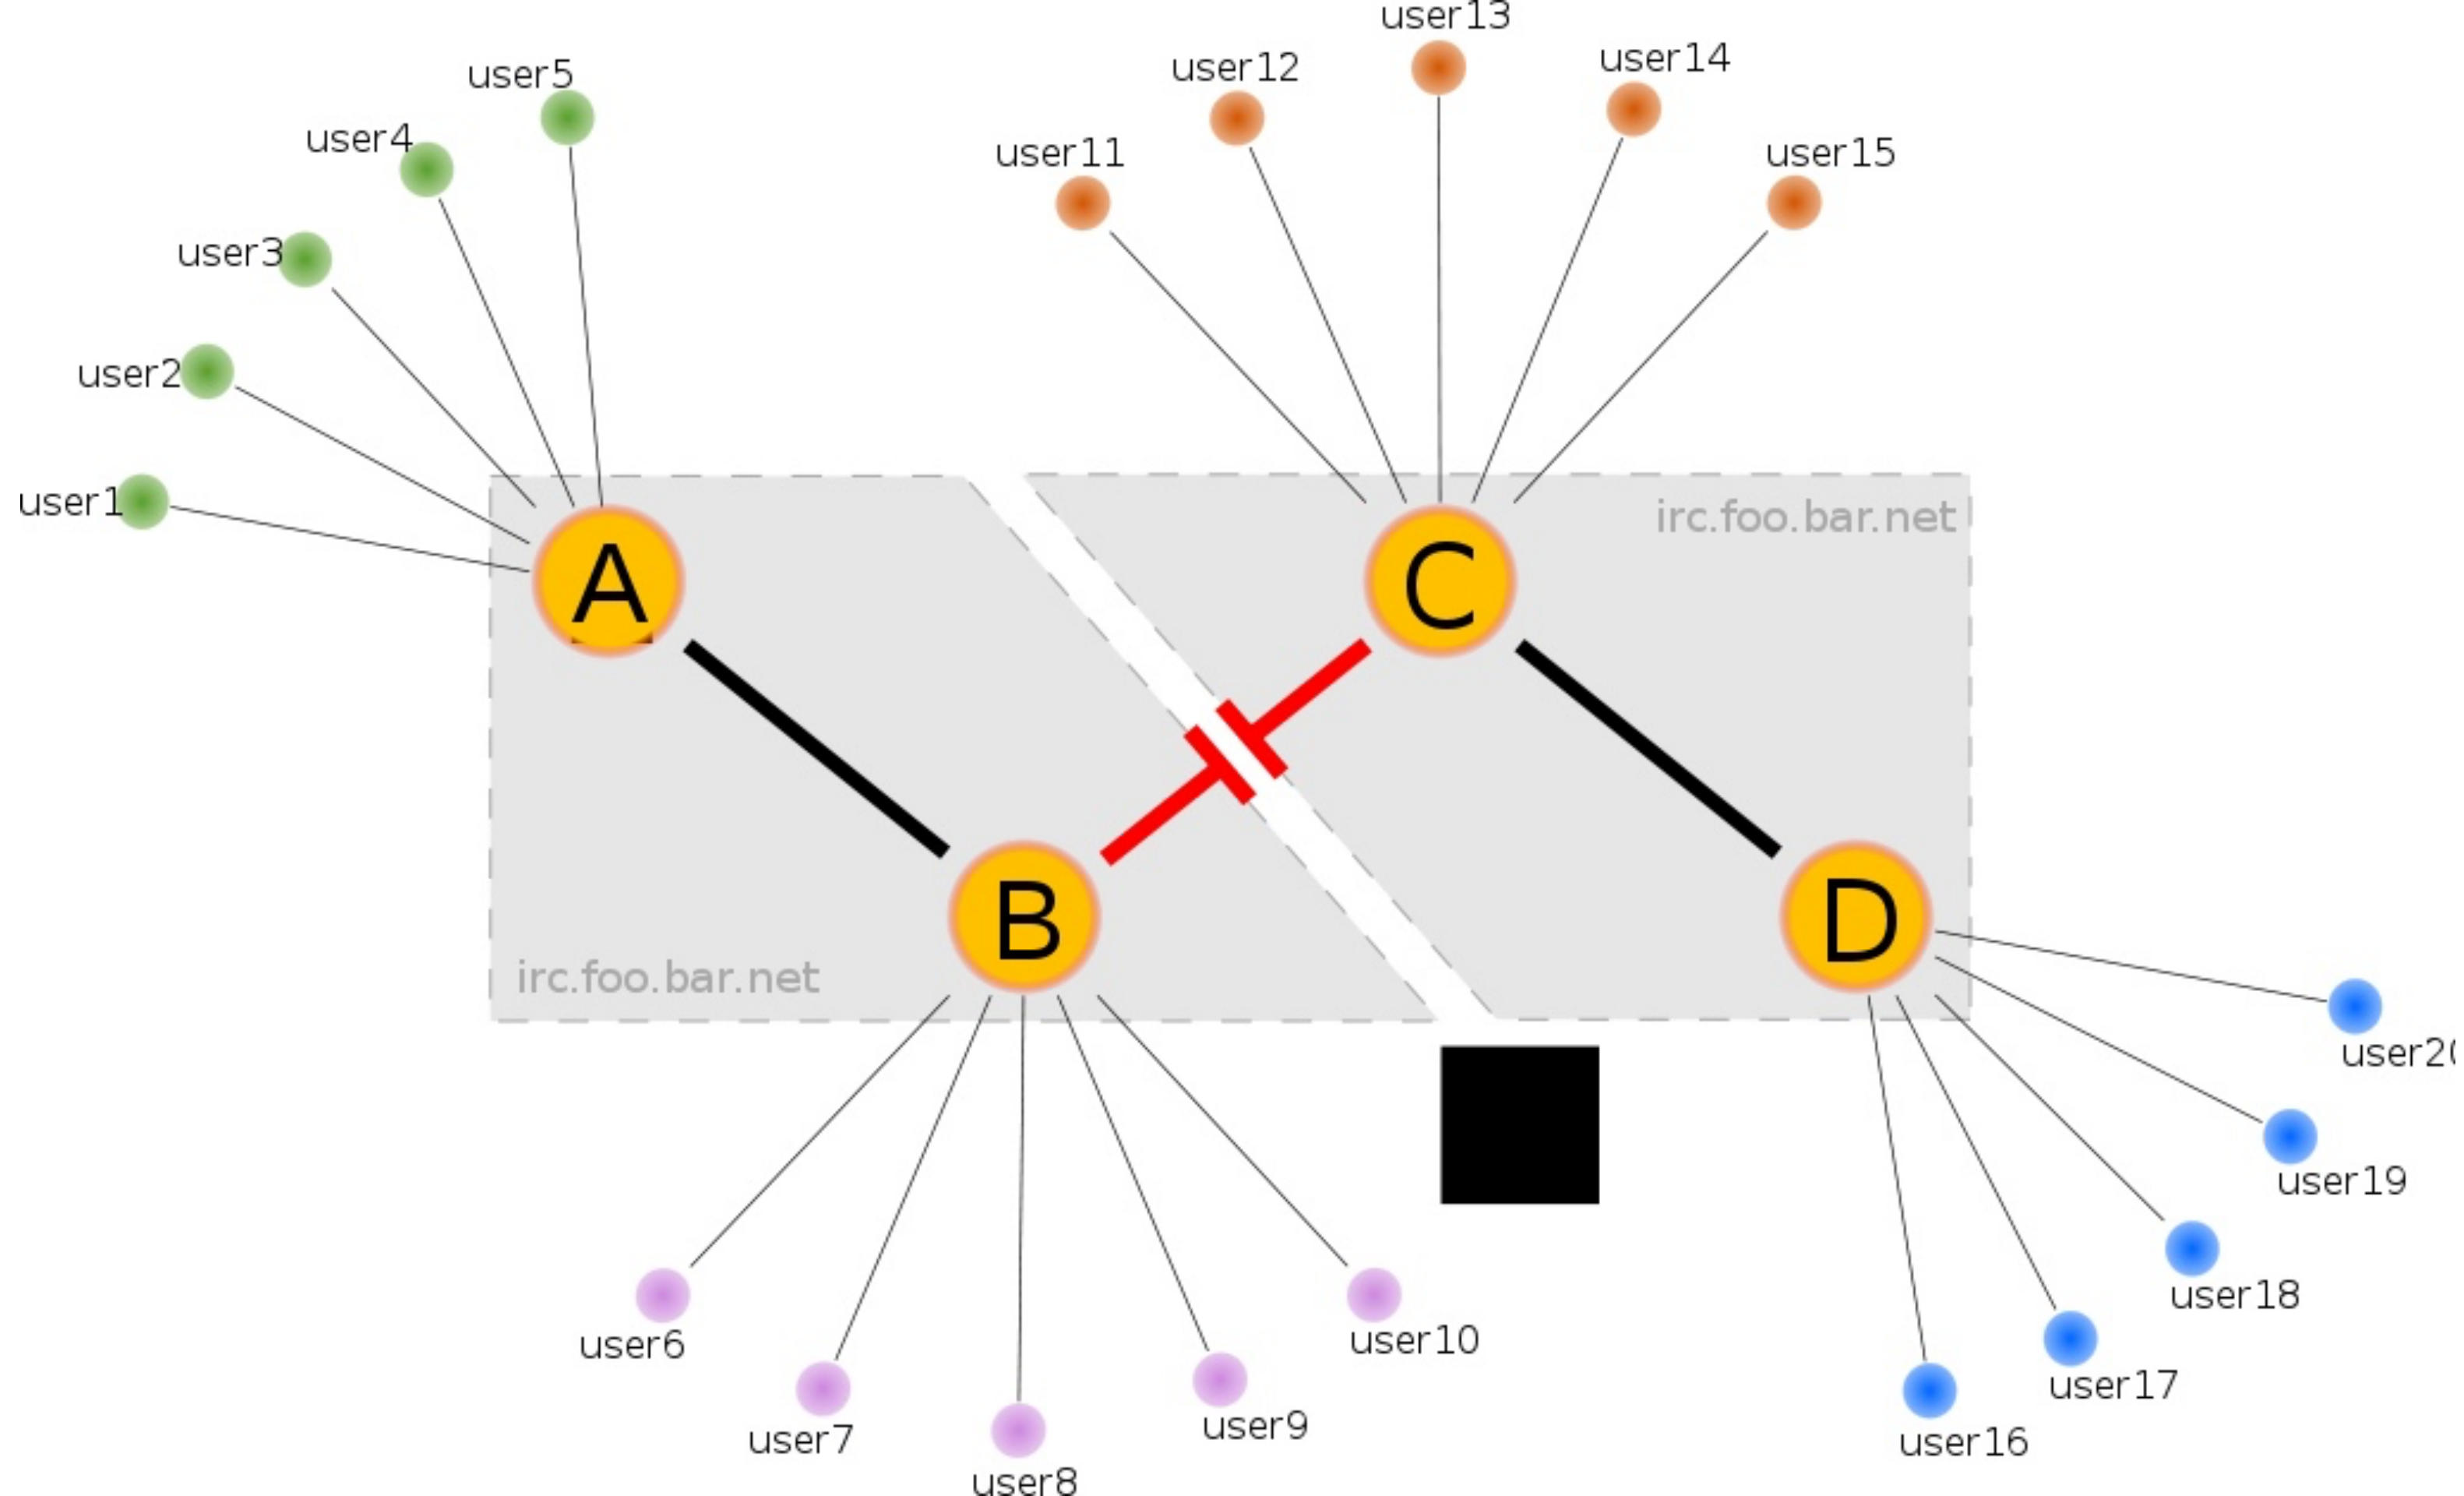
Topic of this video:

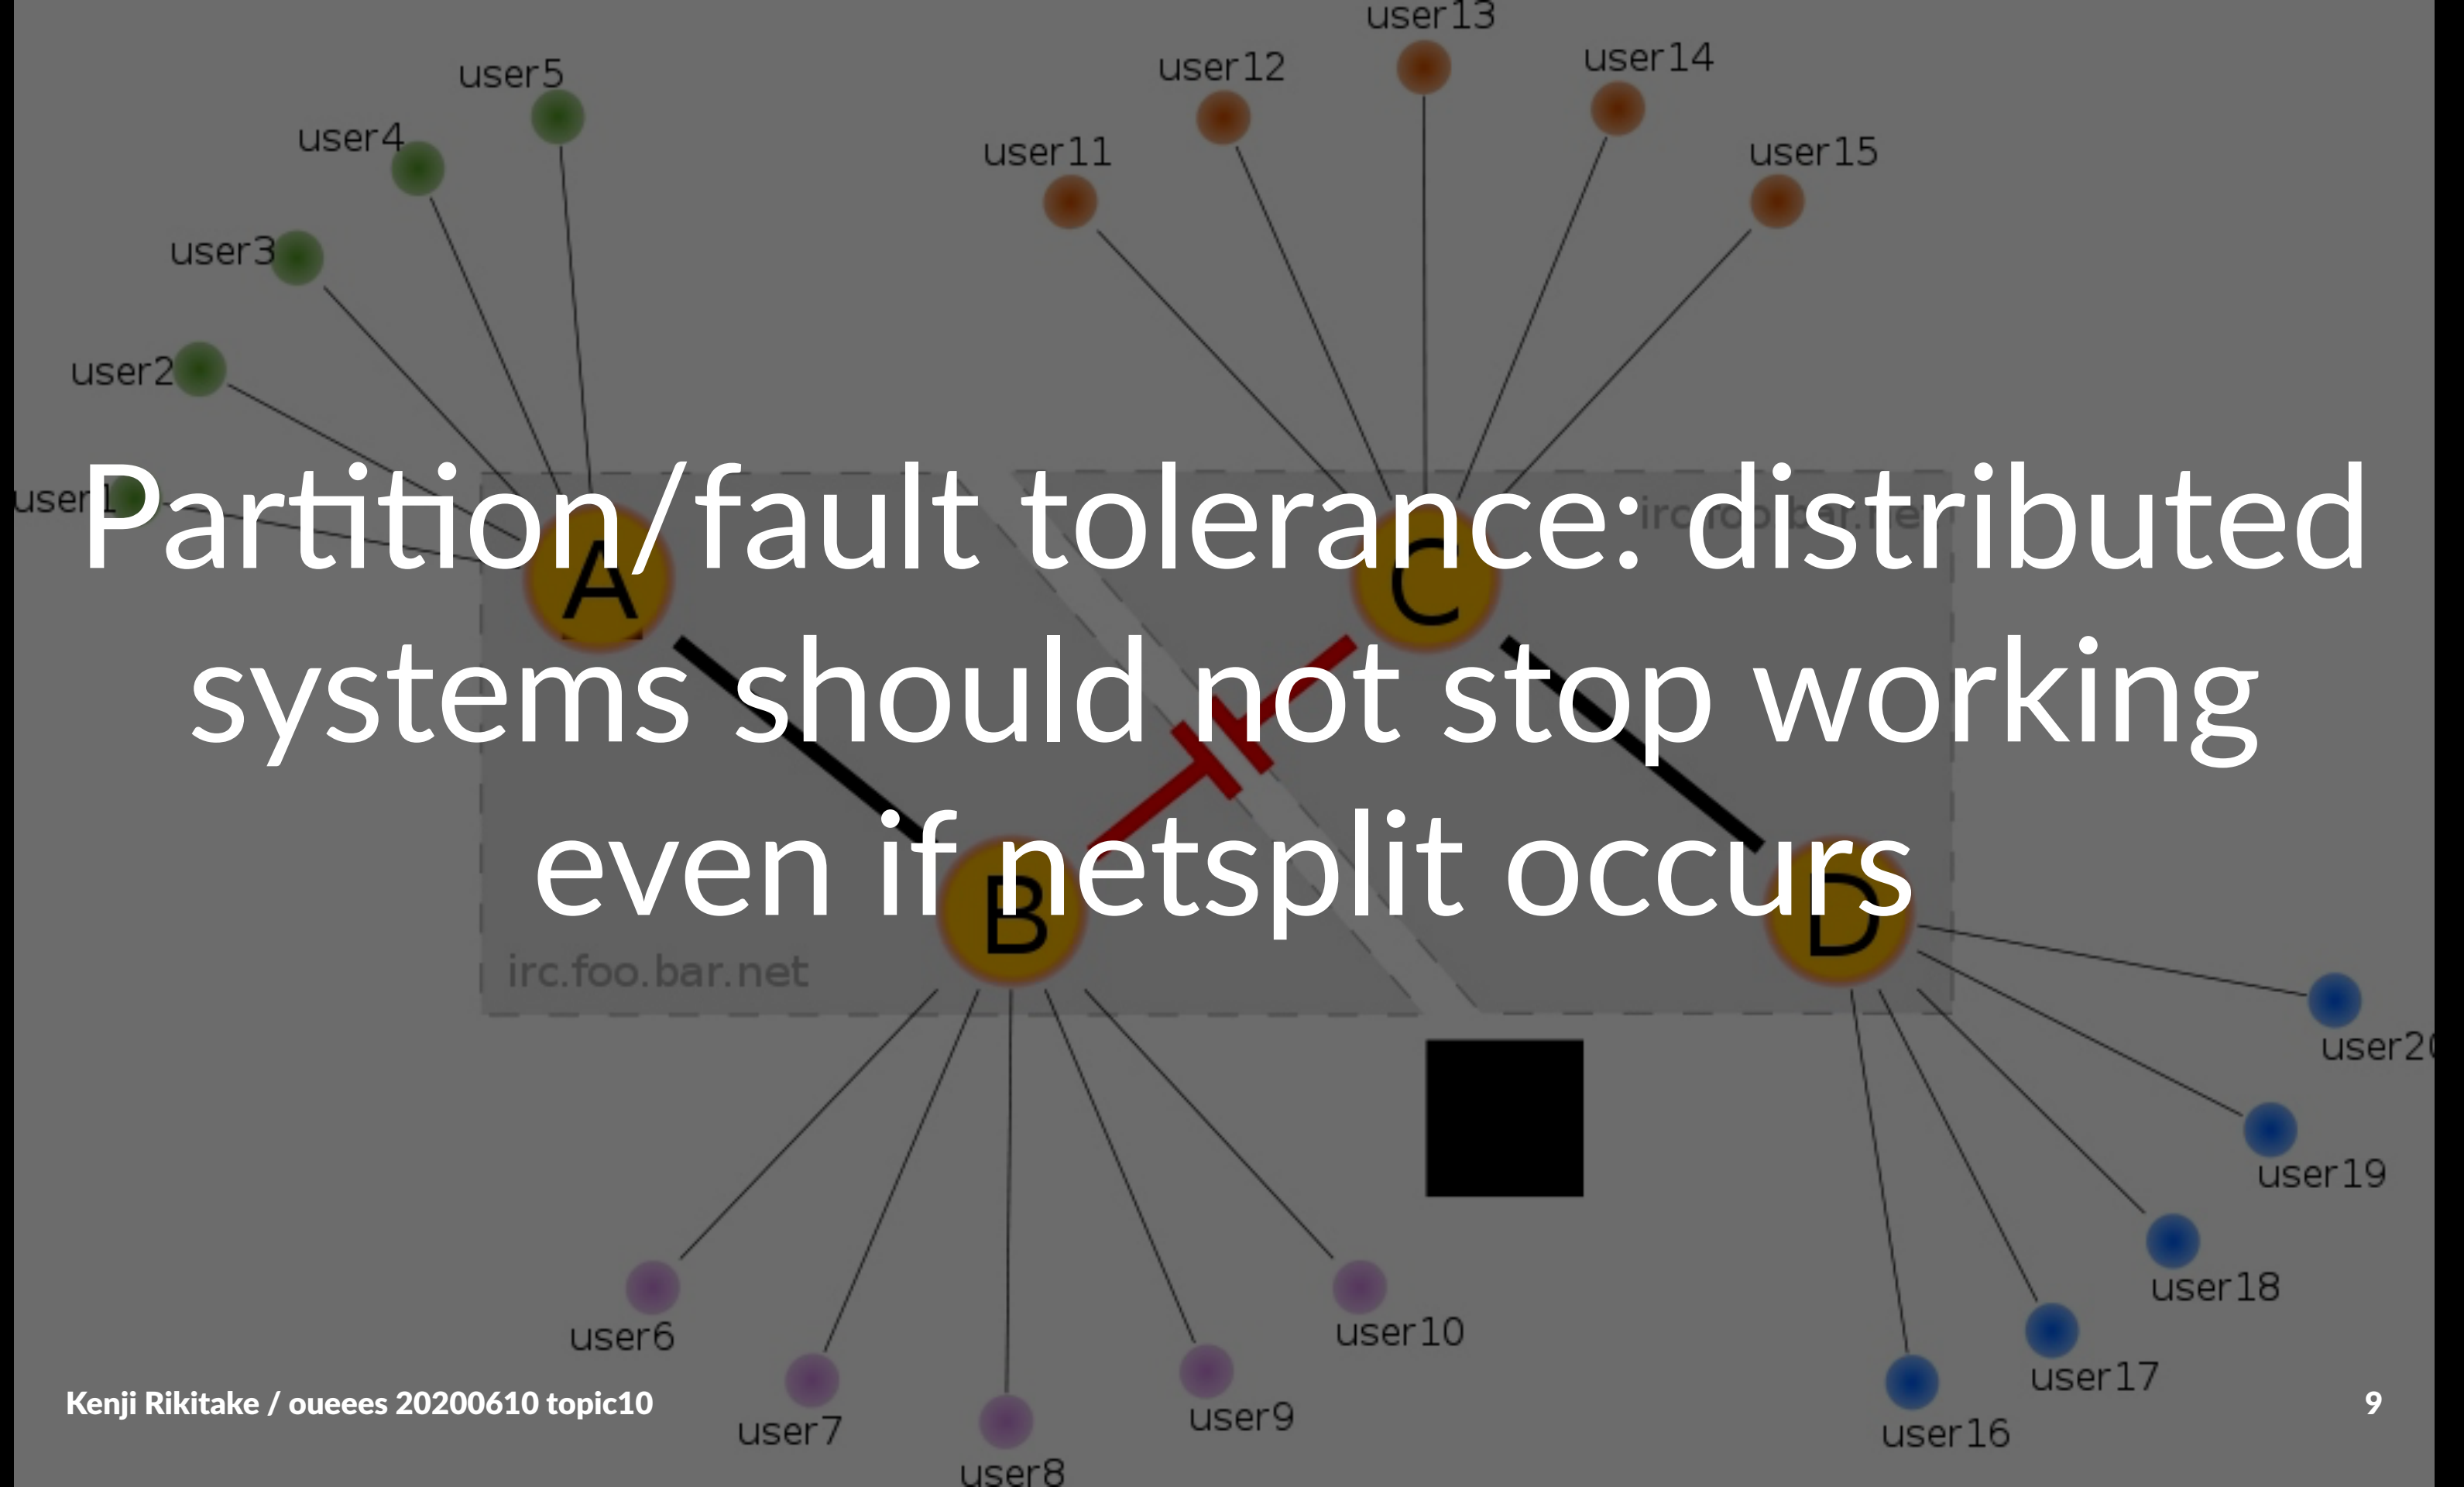
Network fault-tolerance





Networks *split*





Real-world challenges

- Natural disasters
- Device failures
- Human operation errors
- Political impediments
- Social resentments

Handling *failures*

- Redundancy: keeping backup units ready
- Fault tolerance: keeping systems running even the components fail
- Resilience by failing fast: early detection of failures and invocation of the recovery procedures

Why fault tolerance?

- Hard disk MTBF \approx 1 million hours
- 1000 hard disks running 24 hours x 365 days = 8.76 million hours
- If you're running a system with 1000 hard disks, **9 out of 1000** will fail in a year
- Recovery of a disk content takes often *a day*; you can't stop a system for *a day*, can you?

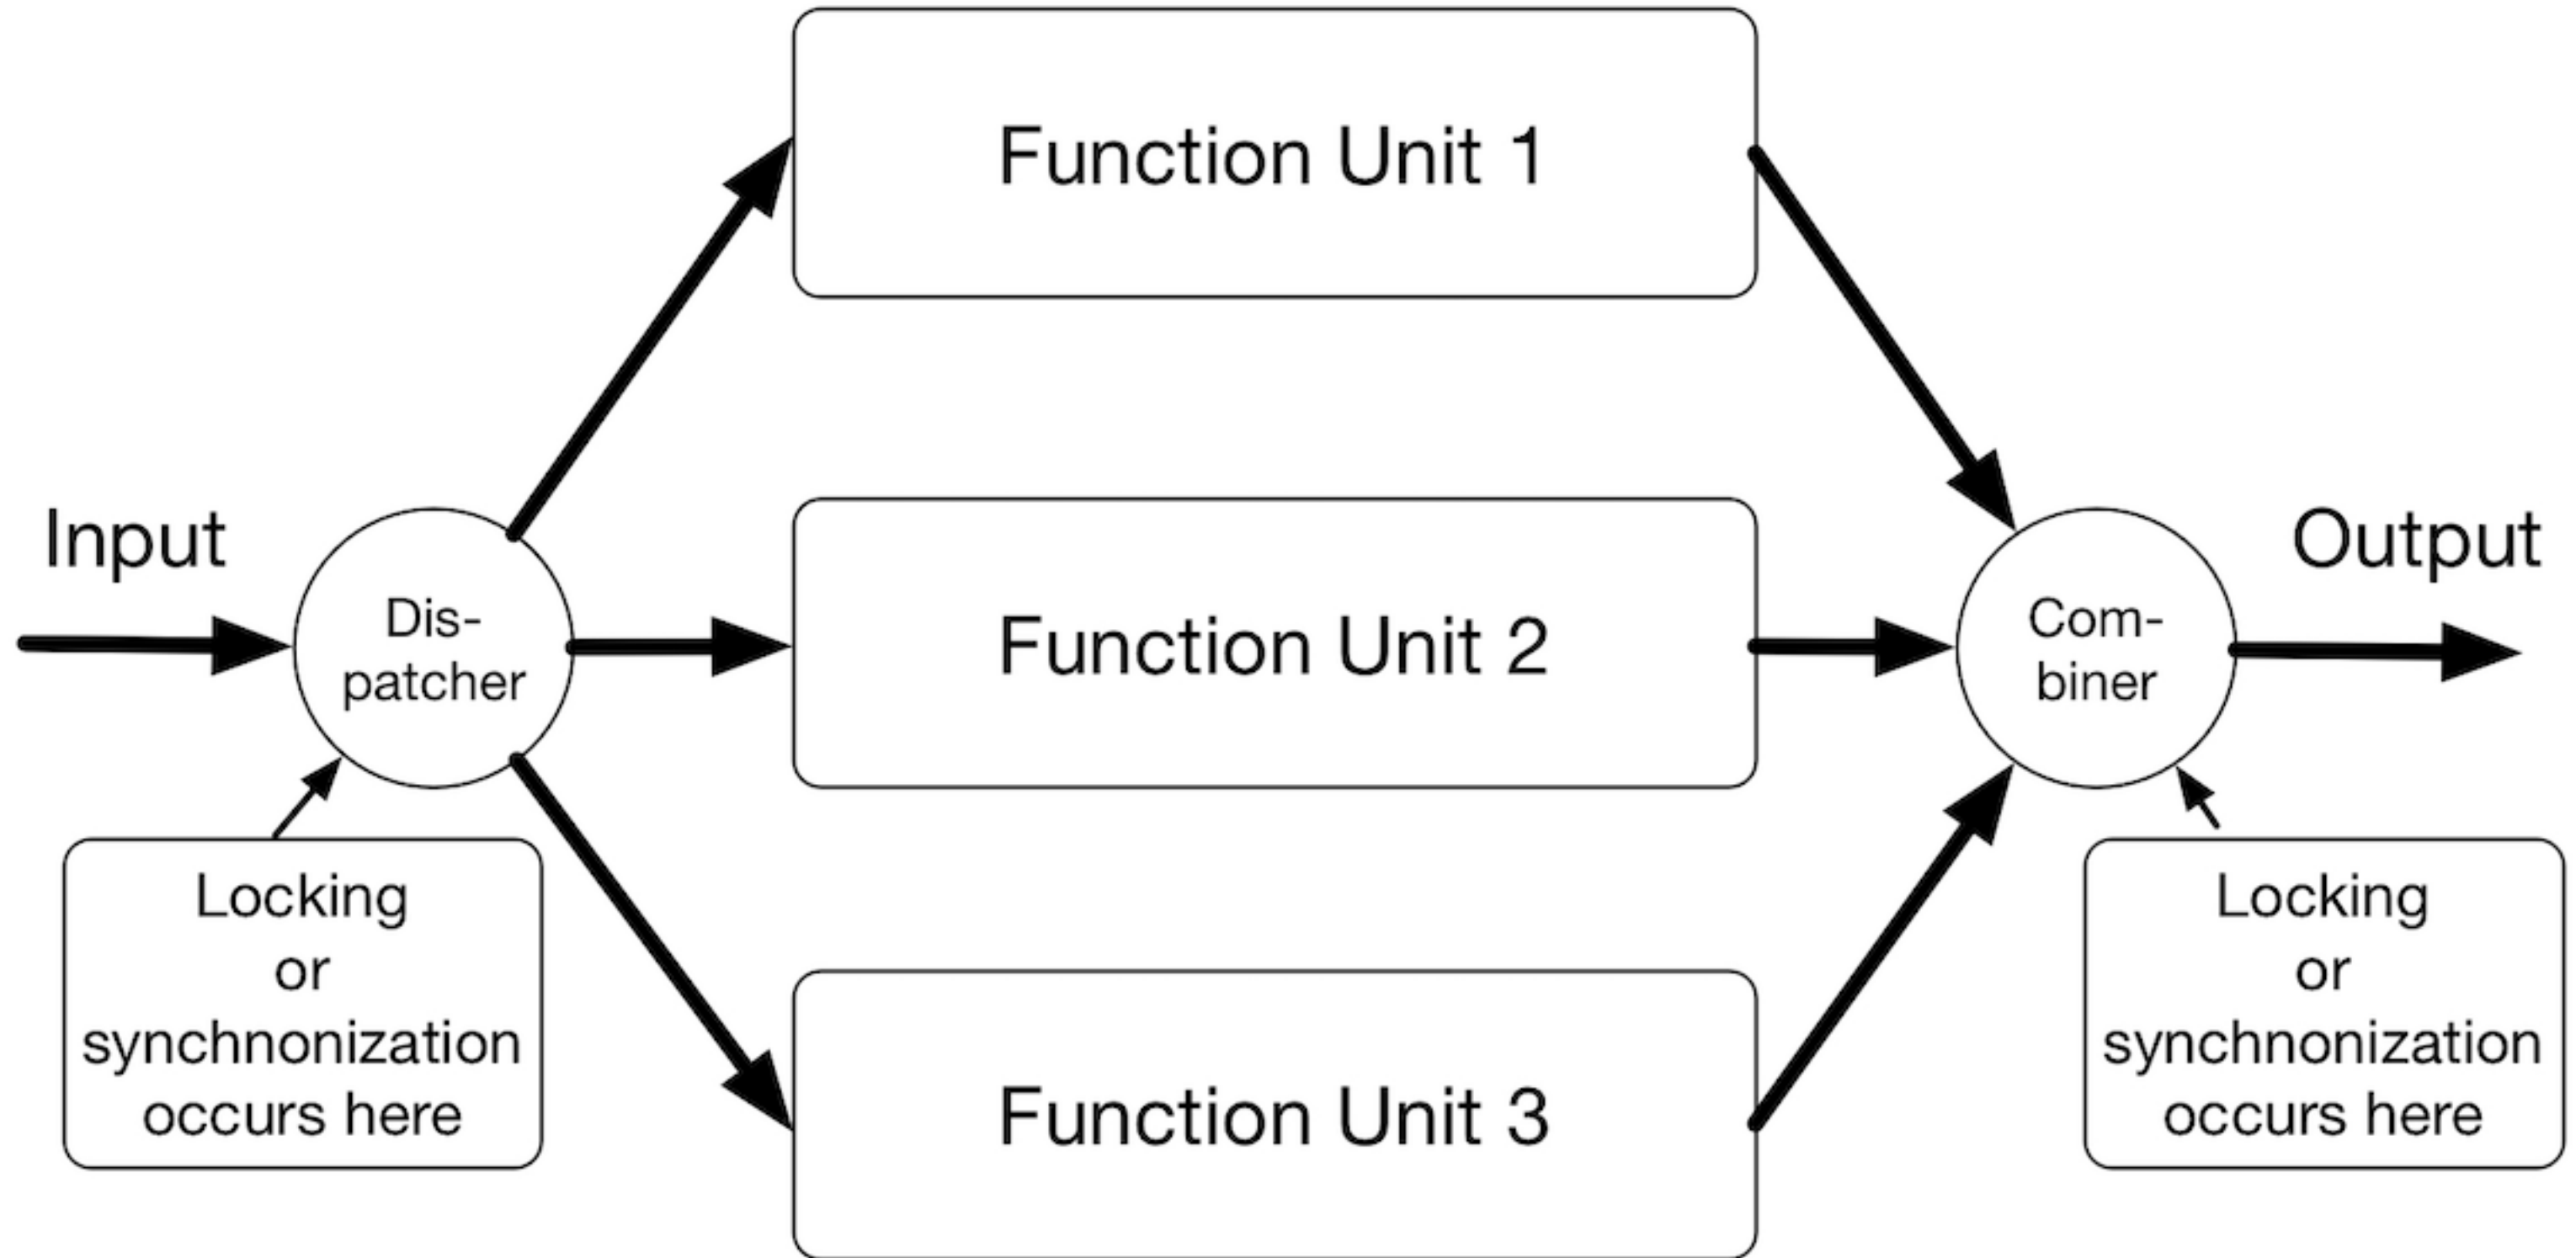
Requirement to keep the systems fault tolerant

- Redundancy: two or more resources for each unit of processing
- Supervising the failure of the units by an independent supervisor
- Rollback capability: undo the incomplete operations and retry

Consistency issues of distributed systems


- Locking/synchronization: waiting all data to be ready to compute or proceed to next step
- Choosing the *right* data: which data is *correct*?
- Supervision: fault detection and restarting

Each function unit runs on
its own speed



Supervision
tree example

Try to
restart



when
crashed

Eight Fallacies of Distributed Computing³ (1/2)

- **The network is reliable**
- **Latency is zero**
- **Bandwidth is infinite**
- The network is secure

³ <https://blog.fogcreek.com/eight-fallacies-of-distributed-computing-tech-talk/>

Eight Fallacies of Distributed Computing (2/2)

- Topology doesn't change
- There is one administrator
- Transport cost is zero
- The network is homogeneous

Summary: centralized computing is
fragile; distributed computing is
fault tolerant but hard

Photo and image credits

- All photos and images are modified and edited by Kenji Rikitake
- Photos are from Unsplash.com unless otherwise noted
- Networks: Irina Blok
- Networks Split: Pietro De Grandi
- Netsplit: https://commons.wikimedia.org/wiki/File:Netsplit_split.svg, in public domain